

# Introduction to the iPOP-UP HPC cluster

Julien Rey   Olivier Kirsh   Magali Hennion

Juin 2023



# Table of Contents



# Who is this training for

- You work at Université Paris Cité
- You need (or might need) more computational power than you currently have
- You are familiar with Unix systems and Bash






# Missions




- Development of Structural Bioinformatics methods
- Services deployment on the computing resource
- Expertise and training
- Hardware hosting

# Web portal

- <https://mobyle.rpbs.univ-paris-diderot.fr>: free access
- 30+ Structural Bioinformatics services
- 1 million CPU hours /year
- 80,000 jobs /year



# The iPOP-UP project



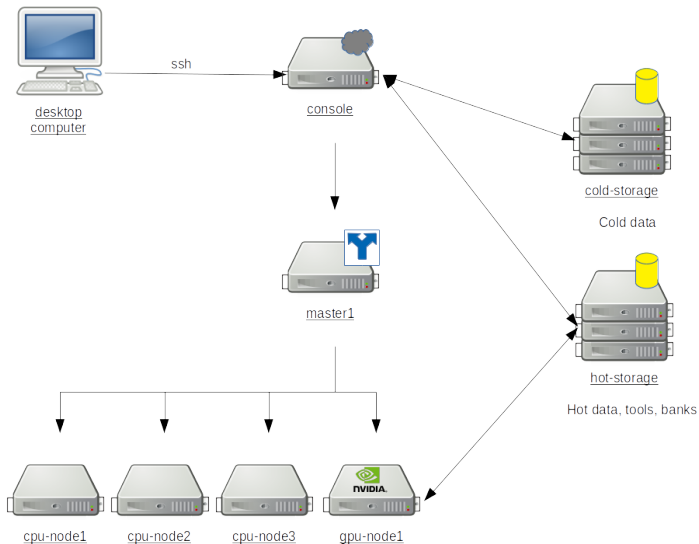
- iPOP-UP: Integrative Platform for Omics Projects at Université de Paris
- Multisite Bioinformatics platform
- Ranging from multiple 'Omics' techniques to structural and chemo- *in silico* Bioinformatics
- Compute nodes

# What is a cluster for ?

- High hardware resources needs
- Long running analyses
- A lot of similar analyses
- Shared work between users
- Free your desktop from the task



# What is a cluster ?





# Computational hardware

Partitions (groups of compute nodes):

- ipop-up: 16 nodes, 2048 CPUs
- rpbs: 20 nodes, 832 CPUs, 9 GPUs (+ 9 nodes, 520 CPUs, 4 GPUs)
- cmpli: 6 nodes, 264 CPUs, 14 GPUs (+ 1 node, 64 CPUs, 2 GPUs)
- epigen: 4 nodes, 128 CPUs
- master-bi: 1 node, 32 CPUs, 3 GPUs

Storage:

- hot-storage: 125TB, very fast
- cold-storage: 240TB, slow + backup (soon)



# In practice

Go to your terminal and connect to the cluster using the following line, don't forget to replace username with your personal username.

```
ssh username@ipop-up.rpbs.univ-paris-diderot.fr
```

Type in your password and enter.

```
You@YourComputer:~/PathTo/RNAseqProject$ ssh username@ipop-up.rpbs.univ-paris-diderot.fr
#####
#
#      iPOP-UP
#
# Hosted by:
# Ressource Parisienne en Bioinformatique Structurale
# -----
# All connections are monitored and recorded.
# Disconnect IMMEDIATELY if you are not an authorized user!
#
#####
username@ipop-up.rpbs.univ-paris-diderot.fr's password:
Last login: Tue Jan 25 15:21:45 2022 from 172.28.18.162
Bienvenue sur le cluster iPOP-UP.

Pour toute question ou demande de support, rejoignez-nous sur le forum de RPBS : https://discourse.rpbs.univ-paris-diderot.fr

Pour changer le compte projet par défaut : sacctmgr update user $USER set defaultaccount=<project-name>
project1 [-----] 3 / 1024 GB
project2 [#-----] 913 / 10240 GB
Update: 2022-01-21 16:00 - default account in bold - More info: status_bars --help
[username@ipop-up ~]$
```



## Note about ssh and security

Your IP will be banned after 5 failed authentication attempts. Each public key count as one authentication attempt. To disable public key authentication:

```
ssh -o PubkeyAuthentication=no  
↪ username@ipop-up.rpbs.univ-paris-diderot.fr
```

A good thing to do, change your password:

```
[rey@ipop-up ~]$ passwd
```

# Where you can go, write, or execute

## User environments

/shared/home/username

## Computations (hot data)

/shared/projects/projectname

## Processed data (cold data)

/cold-storage/username

## Data banks (read-only)

/shared/banks/

# Slurm



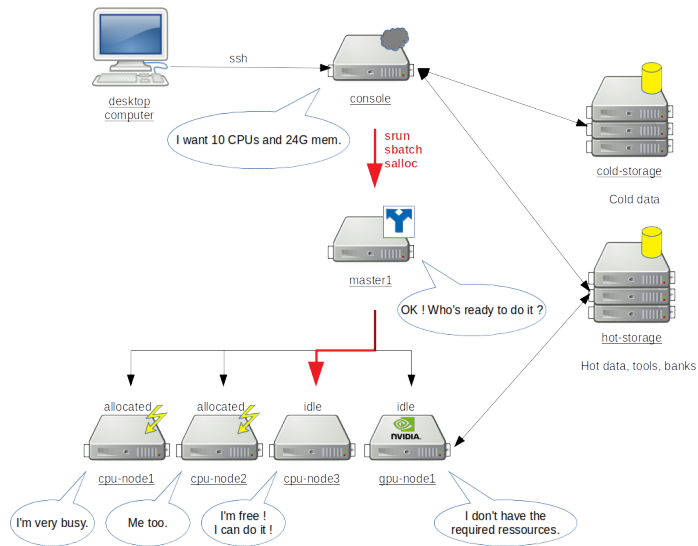
Slurm is the job scheduling system.

It is what will take your code and distribute it on the computing nodes, while ensuring they have the CPU(s) and RAM that you asked for.

It requires specific commands to run (srun, sbatch, salloc, etc...)



# Slurm



# srun

Launch a (simple) interactive job.

```
[rey@ipop-up ~] srun hostname
```

Some parameters can be added to the command line:

- partition**/–**p**: request a specific partition
- account**/–**A**: select the (project) account
- cpus-per-task**/–**c**: request that ncpus be allocated (default: 1 cpu)
- mem-per-cpu**: specify the required memory per cpu (default: 2GB)

Example:

```
[rey@ipop-up ~] srun -A training -p ipop-up -c 8 hostname
```



# sbatch

Launch more complex jobs.

myscript.sbatch

```
#!/bin/bash
#SBATCH --partition=ipop-up
#SBATCH --account=training
#SBATCH --cpus-per-task=8
#SBATCH --mem-per-cpu=4GB
#SBATCH --output=resultat.log
hostname
```

Example:

```
[rey@ipop-up ~] sbatch myscript.sbatch
```



# Associations

Cluster, account, partition and qos must match your user's associations.  
To check your associations:

```
[rey@ipop-up ~] sacctmgr show user rey withassoc
```

User	Def Acct	Cluster	Account	Partition	QOS	Def QOS
rey	demo	production	training	ipop-up	normal	
rey	demo	production	demo	ipop-up	normal	
rey	demo	production	alphafold	cmpli	normal	
rey	demo	production	alphafold	rpbs	normal	

# squeue

List submitted jobs on the cluster:

```
[rey@ipop-up ~] squeue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
1008002	cmpli	argtoser	domingue	R	20:58:09	1	gpu-node14
993325	cmpli	calcam-p	ghoula	R	6-20:23:17	1	gpu-node8
993574	ipop-up	amyloid_	meuret	PD	0:00	1	(AssocGrpCPUMinutesLimit)
943019_[165-299]	ipop-up	kappa_va	badaoui	PD	0:00	1	(JobArrayTaskLimit)
943019_117	ipop-up	kappa_va	badaoui	R	7-17:55:27	1	cpu-node132

Some parameters can be added to filter jobs or show more infos:

- partition/-p**: specify the partition to view
- user/-u**: request jobs from a list of users
- format/-o**: specify the information to be displayed

Example:

```
[rey@ipop-up ~] squeue --me -p ipop-up
[rey@ipop-up ~] squeue -o "%10i %.9P %.8j %.12u %.15T %.15M %.15l %.6D %.20R %.6C %.10b %.10Q"
```



# scancel

Kill a job:

```
[rey@ipop-up ~] scancel job_id
```

Display accounting data for your running and completed jobs:

```
[rey@ipop-up ~] sacct --format=JobID,JobName,Start,Elapsed,NCPUS,NodeList,ReqMeM,State --start 2023-06-01 --en
```

JobID	JobName	Start	Elapsed	CPUTime	NCPUS	NodeList	ReqMem	State
1010778	sleep	2023-06-08T15:37:19	00:00:30	00:00:30	1	cpu-node133	2000Mc	COMPLETED

## Report a job's efficiency:

```
[rey@ipop-up ~]$ seff 981654
Job ID: 981654
Cluster: production
User/Group: lejal/cmpli
State: COMPLETED (exit code 0)
Cores: 1
CPU Utilized: 04:15:58
CPU Efficiency: 99.72% of 04:16:41 core-walltime
Job Wall-clock time: 04:16:41
Memory Utilized: 243.88 GB
Memory Efficiency: 97.01% of 251.40 GB
```

# Some vocabulary

A *job* consists of one or more *steps*, each consisting of one or more *tasks* each using one or more *CPUs*.

- job: A script, typically started with *sbatch*
- step: A step in the job, typically started with *srun*
- task: Requested at the job or step level, with *-array* or *-ntasks*



# Job arrays

Job arrays offer a mechanism for launching a lot of tasks at the same time. Each task of the job will have the environment variable `$SLURM_ARRAY_TASK_ID` set to its array index value.

`myscript.sbatch`

```
#!/bin/bash
#SBATCH --partition=ipop-up
#SBATCH --account=training
#SBATCH --output=resultat_%a.log
#SBATCH --array=1-3
case "$SLURM_ARRAY_TASK_ID" in
  1) fruit='orange';;
  2) fruit='apple';;
  3) fruit='banana';;
esac
echo $fruit
```

# Job arrays

```
[rey@ipop-up ~]$ sbatch myscript.sbatch
```

Results:

```
[rey@ipop-up ~]$ ls
resultat_1.log resultat_2.log resultat_3.log
[rey@ipop-up ~]$ tail resultat_*
==> resultat_1.log <==
orange

==> resultat_2.log <==
apple


==> resultat_3.log <==
banana
```



# Job Array Common Mistakes

- The index of bash lists starts at 0
- Don't forget to have different output files for each task of the array
- Same with your log names (%a or %J in the name will do the trick)
- Do not overload the cluster! Please use %50 (for example) at the end of your indexes to limit the number of tasks (here to 50) running at the same time. The 51st will start as soon as one finishes!
- The RAM defined using #SBATCH --mem=25G is for each task

# Complex workflows



Use workflow managers such as Snakemake or Nextflow.

# Tools

Tools are installed on the cluster in virtual environments:

- each tool has its own dependencies (libraries) and it's not possible to make them all coexist in the same environment
- reproducibility: some users need a specific version of a tool



Conda environments



Containers (Apptainer)

# Modules

They can be loaded with the module command.

Look for the different versions of multiqc:

```
[rey@ipop-up ~]$ module avail multiqc  
multiqc/1.3  multiqc/1.6  multiqc/1.7  multiqc/1.9
```


Load an environment:

```
[rey@ipop-up ~]$ module load multiqc/1.9
```

List loaded environments:

```
[rey@ipop-up ~]$ module list  
Currently Loaded Modulefiles:  
  1) multiqc/1.9      2) blast/2.13.0
```

# How tools are installed



## Useful resources

To find out more, the SLURM manual : `man sbatch` or  
<https://slurm.schedmd.com/sbatch.html>

Ask for help or signal problems on the cluster :  
<https://discourse.rpbs.univ-paris-diderot.fr/>

iPOP-UP cluster documentation:  
<https://ipop-up.docs.rpbs.univ-paris-diderot.fr/documentation/>

# Thanks



- Alix Silvert

iPOP-UP's technical and steering committees





[https://parisepigenetics.github.io/bibs/cluster/training\\_230612/training/](https://parisepigenetics.github.io/bibs/cluster/training_230612/training/)